

# Assignment 3

**Student Name: Krishna Vamsi Koppula**

**Student ID: 700745021**

**GitHub Link:**

[https://github.com/Krishnavamsikoppula/Assignment\\_3\\_NNDL\\_Summer](https://github.com/Krishnavamsikoppula/Assignment_3_NNDL_Summer)

**Video Link:**

<https://drive.google.com/file/d/1KWLwhk4mCklbBsRfarTvNvER5n2FliTo/view?usp=sharing>

1. Follow the instruction below and then report how the performance changed. (Apply all at once)
  - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Flatten layer.
  - Dropout layer at 20%.
  - Fully connected layer with 1024 units and a rectifier activation function.
  - Dropout layer at 20%.
  - Fully connected layer with 512 units and a rectifier activation function.
  - Dropout layer at 20%.

- Fully connected output layer with 10 units and a Softmax activation function



NNDL\_700745021\_ICP3.ipynb ☆

File Edit View Insert Runtime Tools Help [All changes saved](#)

+ Code + Text



{x}



```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.optimizers import SGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
```

✓  
1m

#### # Create the model

```
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = SGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
```

✓  
1m

#### # Evaluate the model

```
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))
```

## Output:

✓ 1m	Model: "sequential_1"		
	Layer (type)	Output Shape	Param #
	conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
	dropout_6 (Dropout)	(None, 32, 32, 32)	0
	conv2d_7 (Conv2D)	(None, 32, 32, 32)	9248
	max_pooling2d_3 (MaxPooling 2D)	(None, 16, 16, 32)	0
	conv2d_8 (Conv2D)	(None, 16, 16, 64)	18496
	dropout_7 (Dropout)	(None, 16, 16, 64)	0
	conv2d_9 (Conv2D)	(None, 16, 16, 64)	36928
	max_pooling2d_4 (MaxPooling 2D)	(None, 8, 8, 64)	0
	conv2d_10 (Conv2D)	(None, 8, 8, 128)	73856
	dropout_8 (Dropout)	(None, 8, 8, 128)	0
	conv2d_11 (Conv2D)	(None, 8, 8, 128)	147584
	max_pooling2d_5 (MaxPooling 2D)	(None, 4, 4, 128)	0
	flatten_1 (Flatten)	(None, 2048)	0
	dropout_9 (Dropout)	(None, 2048)	0
	dense_3 (Dense)	(None, 1024)	2098176
	dropout_10 (Dropout)	(None, 1024)	0
	dense_4 (Dense)	(None, 512)	524800
	dropout_11 (Dropout)	(None, 512)	0
	dense_5 (Dense)	(None, 10)	5130
	Total params: 2,915,114 Trainable params: 2,915,114 Non-trainable params: 0		
	None		
	Epoch 1/5		
	1563/1563 [=====] - 14s 8ms/step - loss: 1.8837 - accuracy: 0.3061 - val_loss: 1.5697 - val_accuracy: 0.4337		
	Epoch 2/5		
	1563/1563 [=====] - 12s 8ms/step - loss: 1.5043 - accuracy: 0.4499 - val_loss: 1.4281 - val_accuracy: 0.4964		
	Epoch 3/5		
	1563/1563 [=====] - 12s 8ms/step - loss: 1.3666 - accuracy: 0.5030 - val_loss: 1.3071 - val_accuracy: 0.5325		
	Epoch 4/5		
	1563/1563 [=====] - 12s 8ms/step - loss: 1.2885 - accuracy: 0.5359 - val_loss: 1.2516 - val_accuracy: 0.5507		
	Epoch 5/5		
	1563/1563 [=====] - 12s 8ms/step - loss: 1.2268 - accuracy: 0.5556 - val_loss: 1.1757 - val_accuracy: 0.5754		
	Accuracy: 57.54%		

Did the performance change?

With the addition of more layers and feature maps, the model's performance is likely to improve, but the complexity and training time of the model will also rise. The new model architecture described in the instruction has more feature maps and new layers, which could increase the model's accuracy.

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.

✓  
0s



```
# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])

# Convert the predictions to class labels
predicted_labels = np.argmax(predictions, axis=1)

# Convert the actual labels to class labels
actual_labels = np.argmax(y_test[:4], axis=1)

# Print the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels:", actual_labels)
```



```
1/1 [=====] - 0s 324ms/step
Predicted labels: [3 1 8 0]
Actual labels: [3 8 8 0]
```

### 3. Visualize Loss and Accuracy using the history object.

```
import matplotlib.pyplot as plt

# Plot the training and validation loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

# Plot the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

